

Nymish Kashivishwanath

reachnymish@gmail.com | linkedin.com/in/nymish | github.com/nymishkash | Bengaluru, India

EXPERIENCE

AI Studio

Jun 2026 – Present

Forward Deployed Engineer

Bengaluru, India (Remote)

- Forward-deployed engineer embedding with enterprise clients to scope, build, and ship custom **AI-agent solutions** end-to-end — from discovery to production deployment.
- Build production **LLM systems** (RAG, function-calling agents, workflow automation) integrated directly into client stacks and data sources.
- Own delivery end-to-end: translate ambiguous business requirements into working systems and iterate on-site with stakeholders.

Wing Assistant (M32 Labs)

May 2025 – Jun 2026

Software Engineer

Berkeley, CA (Remote)

- Engineering hire #3 on an AI executive-assistant product; shipped full-stack features end-to-end on Node.js + React.
- Built a **distributed message-queue ingestion pipeline** (RabbitMQ quorum queues, publisher confirms, Redis idempotent dedupe, DLQ + retries) consuming Gmail and Outlook push notifications with zero dropped events under load.
- Architected a **multi-tier cache** (in-memory → Redis → source API) with Redis Pub/Sub cross-instance invalidation and circuit breakers, cutting p95 read latency from **>10s to <500ms**.
- Built the **queue consumer layer**: a worker fan-out running AI draft generation, classification, task creation, and contact extraction concurrently per event, with shared context to eliminate N+1 queries.
- Built the **AI generation pipeline**: GPT-4 drafts conditioned on user writing style, RAG over a Pinecone knowledge base, and CRM/calendar context injection; migrated the LLM provider live for **~20x cost reduction**.
- Delivered full **Microsoft Graph parity** (OAuth, webhook subscriptions, refresh-token worker) for an originally Google-only product, unblocking the enterprise segment.
- Built an **SSE real-time sync layer** over Redis Pub/Sub with reconnect, event dedup, and mutation cooldowns to suppress stale upstream events during eventual-consistency windows.

Dukaan

Feb 2025 – Apr 2025

Full Stack Engineer

Bengaluru, India (Onsite)

- Core team member (6 engineers) on **bot9.ai**, a GenAI customer-support suite; led engineering in the final months and independently owned enterprise client relationships end-to-end.
- Shipped GenAI automation flows that resolved **80–85%** of inbound support queries with no human handoff.
- Cut LLM infrastructure cost **20x** by migrating providers live with zero downtime and no quality regression.
- Replaced WebSockets with **SSE**, improving streaming reliability and concurrent-session scalability for live chat.

Dukaan

Aug 2024 – Jan 2025

SDE Intern

Bengaluru, India (Onsite)

- Shipped production features across Next.js, Express, and OpenAI APIs; chatbot work fed directly into the full-time role.

SKILLS

Languages: TypeScript, JavaScript, Python, SQL

Backend: Node.js, Express, RabbitMQ, Redis (Pub/Sub, Lua), PostgreSQL, BullMQ, Docker, PM2

Frontend: React 19, TanStack Query, Zustand, Tailwind CSS, Vite

AI / Data: OpenAI (GPT-4 / 4o), Pinecone, RAG, function calling

Integrations: Gmail API, Google Cloud Pub/Sub, Microsoft Graph, OAuth2, HubSpot, Pipedrive, Salesforce, AWS S3

PROJECTS

WageWizard Payroll Assistant

- AI payroll assistant on Node.js + Next.js handling HR workflows (employees, leave, wages) via chat and voice.

EDUCATION

Scaler School of Technology

2023 – 2027

Integrated BS & MS in Computer Science

India

Birla Institute of Technology and Science, Pilani

2023 – 2026

Bachelor of Science in Computer Science

India

ACHIEVEMENTS

Infosys Catch Them Young Scholar (2019); Alumnus, **Junior Science Academy**, New York Academy of Sciences.